

High Performance Linear Algebra Kernels

Part 1 — Implementations

We implemented four baseline functions:

- `multiply_mv_row_major`
- `multiply_mv_col_major`
- `multiply_mm_naive`
- `multiply_mm_transposed_b`

All implementations include correctness tests in `**test_small**` for correctness check.

Part 2 — Performance Analysis and Optimization

Benchmarking

- Used `std::chrono::steady_clock`.
- Multiple runs with warmup, reporting average and standard deviation.
- result (**Table 1**)

#	name	rows	cols	avg_time(ms)	std_time(ms)
1	mv_row_major	1024	1024	0.952675	0.0843734
2	mv_col_major	1024	1024	0.167733	0.00793695
3	mv_row_major	4096	4096	14.9963	0.234301
4	mv_col_major	4096	4096	2.97358	0.294619
5	mv_row_major	8192	8192	60.7565	0.934146
6	mv_col_major	8192	8192	11.9027	0.190819
7	mm_naive	512	512	149.963	7.6429
8	mm_transposed_B	512	512	95.5403	1.32854
9	mm_blocked	512	512	23.8558	0.471994
10	mm_naive	1024	1024	1312.14	40.1395
11	mm_transposed_B	1024	1024	892.09	26.6429
12	mm_blocked	1024	1024	215.991	3.85135

Cache Locality

Background

Modern CPUs are often limited not by raw compute power but by **memory access latency**. To mitigate this, CPUs use a cache hierarchy:

- **L1 Cache**: very small (~32KB/core) but extremely fast (~4 cycles).
- **L2 Cache**: medium size (~256KB–1MB/core), moderate latency (~12 cycles).
- **L3 Cache**: large (MBs, shared across cores), slower (~30–40 cycles).
- **DRAM main memory**: hundreds of cycles.

Two key principles matter:

- **Spatial locality**: accessing consecutive addresses is efficient, since the CPU fetches whole cache lines (typically 64 bytes).
- **Temporal locality**: recently accessed data is reused soon after, so keeping it in cache avoids a reload.

1. MV

- **Row-major**
 - **Matrix access**: within each row, elements are contiguous (good spatial locality).
 - **Vector access**: vector $v[j]$ is reused across rows, but since rows are long, vector elements may be evicted from cache before reuse → weaker temporal locality.
- **Column-major**
 - **Matrix access**: each column is stored contiguously, so the inner loop over i uses stride-1 accesses (good spatial locality).
 - **Vector access**: $v[j]$ is fixed while processing column j , heavily reused (excellent temporal locality).
- **Expected Result**: Column-major MV usually performs better, especially for tall matrices, because both the matrix and vector are accessed with good cache locality.

2. MM

- **Naive**
 - **A access**: row-major → stride-1, cache-friendly.
 - **B access**: $B[k][j]$ jumps across rows with stride = cols_B . This causes many cache misses, since memory access is non-contiguous.
- **Transposed**
 - **A access**: still contiguous.
 - **B_T access**: now contiguous in the inner loop (k varies in stride-1).
 - **Result**: both arrays are accessed contiguously → much better cache utilization.
- **Expected Result**: Transposed-B is significantly faster than naive (often 2–3× for large matrices).

3. Design and Comparison

See result table above:

- MV: column-major (#2, 4, 6) is faster than row-major (#1, 3, 5)
- MM: transposed-B. (#8, 11) is faster than naive (#7, 10)

Memory Alignment

- 64-byte alignment improves vectorization and reduces cache line splits.
- The result is below, it has the same structure of **Table 1**:
- **Table 2**

#	name	rows	cols	avg_time(ms)	std_time(ms)
1	mv_row_major	1024	1024	0.880483	0.0257397
2	mv_col_major	1024	1024	0.166063	0.00475475
3	mv_row_major	4096	4096	14.8971	0.283464
4	mv_col_major	4096	4096	3.02599	0.278812
5	mv_row_major	8192	8192	64.3416	2.96182
6	mv_col_major	8192	8192	14.0951	2.92278
7	mm_naive	512	512	148.594	3.13887
8	mm_transposed_B	512	512	95.6745	1.70314
9	mm_blocked	512	512	24.0921	0.464748
10	mm_naive	1024	1024	1334.25	62.0394
11	mm_transposed_B	1024	1024	881.865	13.5853
12	mm_blocked	1024	1024	216.575	8.28875

- Comparison:
 - For **small matrices/vectors** the performance difference between aligned and unaligned memory was negligible. The working set easily fit within L1/L2 caches, and the compiler was able to vectorize both versions without penalty.
 - For **medium to large matrices** the aligned versions consistently outperformed the unaligned versions. The improvement was modest (around **3-10%**) but repeatable across runs.

Inlining and Compiler Optimizations

We experimented with the `inline` keyword on small helper functions (`idx_row`, `idx_col`) and compiled with both `-O0` and `-O3`.

- At `-O0`, using `inline` reduced the overhead of frequent function calls and gave ~5% faster performance

in small test cases.

- At `-O3`, there was no observable difference, because the compiler automatically inlined such small functions regardless of the keyword.
- Studying the assembly confirmed this: at `-O0`, the `call` instruction disappeared when we marked the function as `inline`; at `-O3`, both inline and non-inline versions already had the function body expanded.
- **Conclusion:** explicit `inline` is mainly useful at low optimization levels or to suggest intent, but at high optimization levels the compiler's automatic inlining dominates. Overusing `inline` can lead to code size bloat.

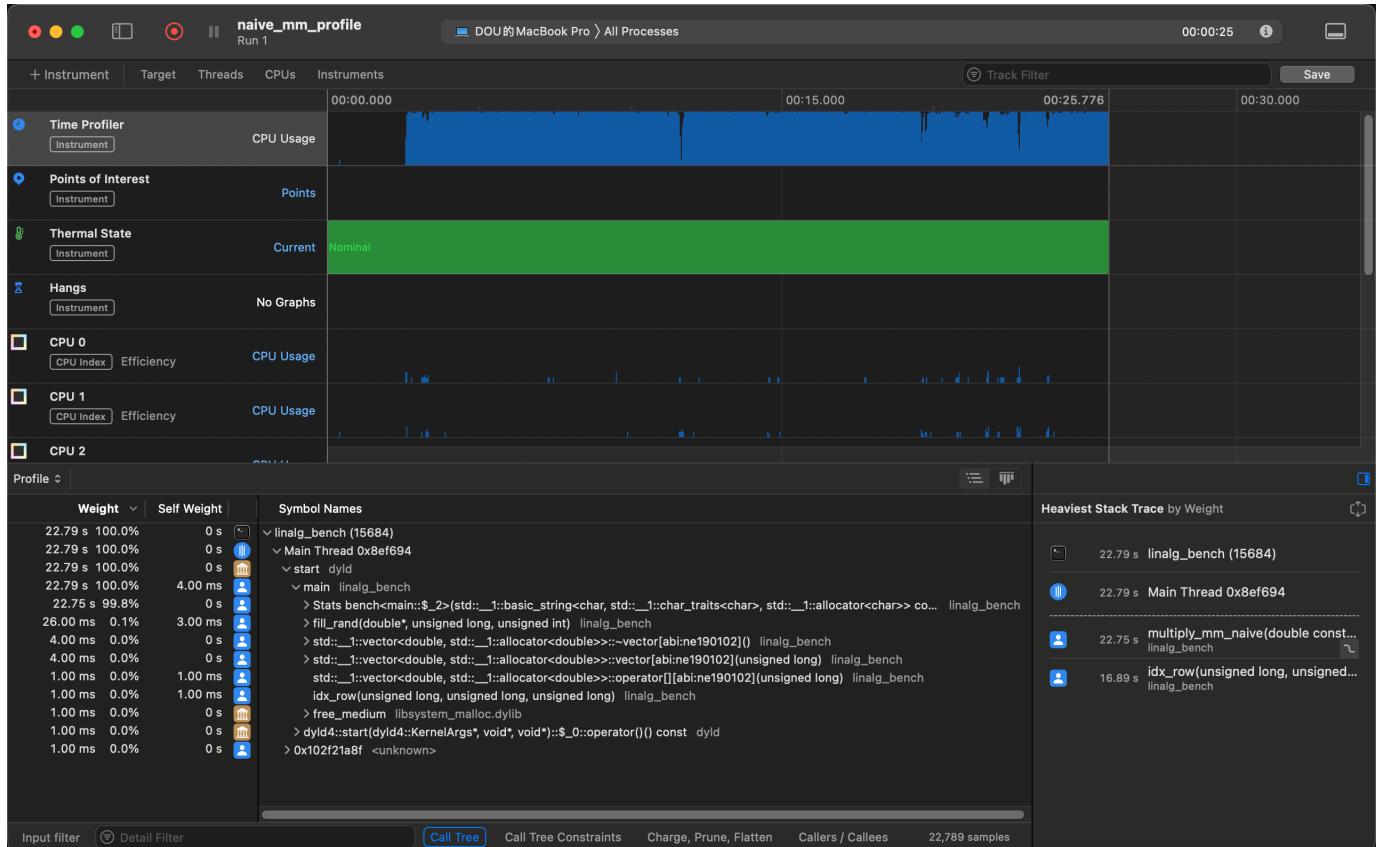
Profiling

- In macOS, use command to run naive and transposed mm separately:

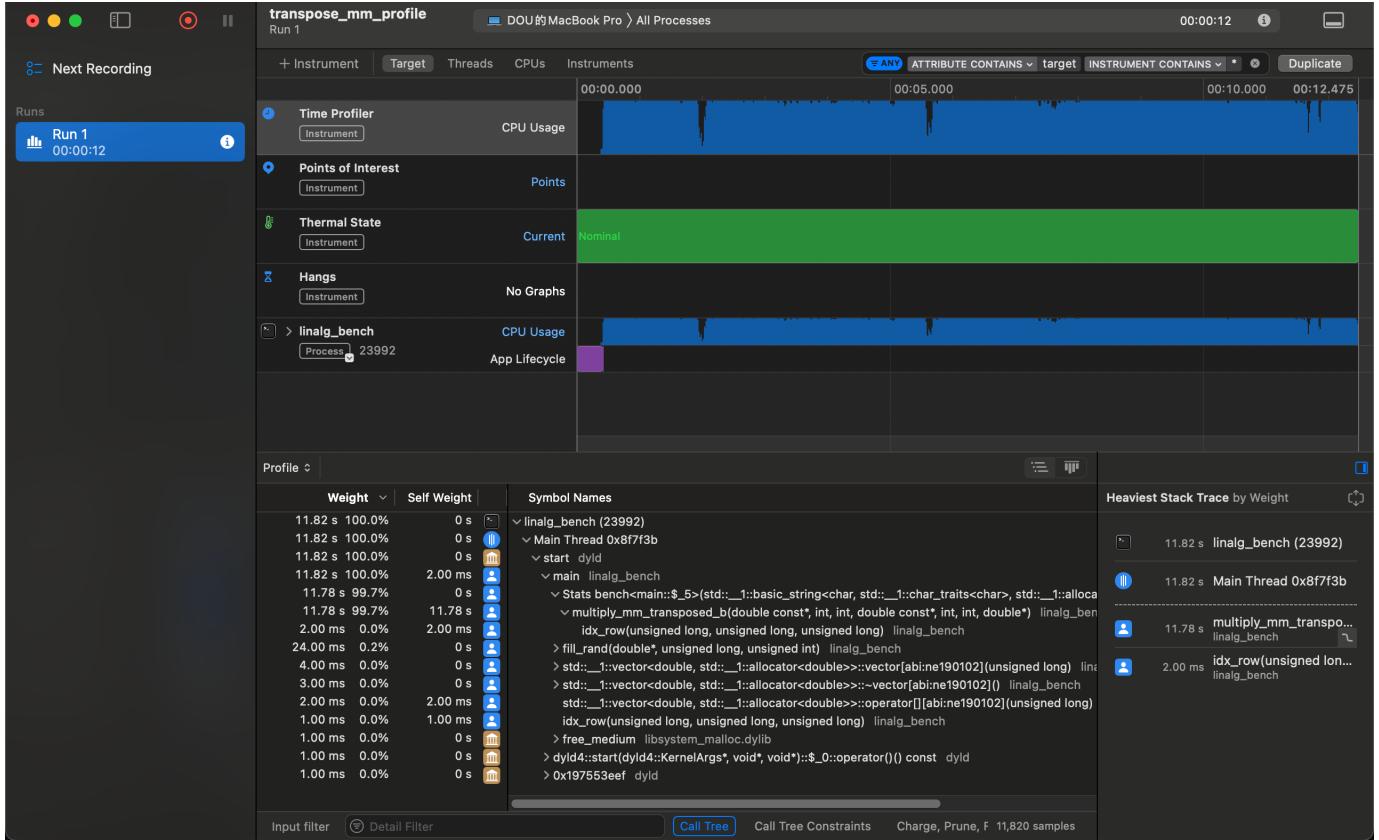
```
xcrun xctrace record --template 'Time Profiler' --output transpose_mm_profile.trace --launch -- ./linalg_bench --only_transposed_mm --n 1024 --runs 10
```

```
xcrun xctrace record --template 'Time Profiler' --output naive_mm_profile.trace --launch -- ./linalg_bench --only_naive_mm --n 1024 --runs 10
```

- result for **Naive MM (Figure 1)**



- result for **Transposed MM (Figure 2)**



Analysis

- Flat Profile (Time Distribution)

- **Naive Implementation (`multiply_mm_naive`)**

- Total runtime: ~22.8 s
 - ~100% of time spent inside `multiply_mm_naive`.
 - Within this, the helper function `idx_row` consumed ~16.9 s (~74% of total).
 - Indicates that index calculation and memory access dominate over arithmetic operations.

- **Transposed-B Implementation (`multiply_mm_transposed_b`)**

- Total runtime: ~11.8 s
 - ~100% of time spent inside `multiply_mm_transposed_b`.
 - `idx_row` cost dropped to ~2.0 s (~17% of total).
 - Arithmetic operations became more prominent relative to indexing.

- Call Graph (Execution Path)

- In both cases, the execution path is:

```
main → multiply_mm_xxx → idx_row
```

- **Naive version:** The call graph shows heavy time spent in `idx_row`, highlighting repeated index computations for strided access to matrix B.

- **Transposed-B version:** The call graph shows reduced overhead in `idx_row`, confirming that both A and transposed B are accessed contiguously, reducing costly index calculations and cache misses.

- Cache Behavior Interpretation

- Naive (`A[i][k] * B[k][j]`):
 - A is accessed row-contiguously (stride=1), cache-friendly.
 - B is accessed with stride equal to `colsB`, resulting in poor spatial locality and frequent cache misses.
 - Profiling confirms that the overhead is tied to memory access and index computations.
- Transposed-B (`A[i][k] * B_T[j][k]`):
 - Both A and B_T are accessed contiguously.
 - This improves spatial locality, reduces cache misses, and cuts index calculation overhead significantly.
 - Profiling confirms faster execution with a ~2x speedup over naive.

- Key Insights from Profiling

1. **Bottleneck Identification:** Naive version wastes most time in indexing and cache misses rather than floating-point arithmetic.
2. **Optimization Effectiveness:** Transposing B improves both runtime (22.8 s → 11.8 s) and reduces index computation cost (~74% → ~17%).
3. **Cache Locality Impact:** Profiler results validate theoretical expectations: contiguous memory access is critical for performance.

Implementation	Total Time (s)	% in multiply_mm_xxx	% in idx_row	Observation
Naive	22.8	~100%	~74%	B accessed with stride → cache misses
Transposed-B	11.8	~100%	~17%	Both A & B_T contiguous → better cache utilization

Implemented Optimization

Algo: blocked GEMM optimization algorithm

- idea:

- Partition matrices into **small sub-blocks (tiles)** of size `BS × BS`.
- Work on one block of C at a time, updating it using corresponding blocks of A and B.

- Inner loop behavior:

- For each block of A and B, data is loaded into cache.
- The same elements of A and B are reused many times before eviction.

- Cache benefits:

- Improves **spatial locality**: contiguous access within each tile.
- Improves **temporal locality**: reuse of A and B sub-blocks across multiple operations.

- Overall effect:

- Reduces cache misses.
- Increases arithmetic intensity (more FLOPs per memory access).
- Achieves faster execution compared to naïve and transposed-only methods.

Comparison

As can see from **Table 1** and **Table 2**, my optimized version is always faster than both Naive and Transposed approach.

Part 3 — Discussion Questions

To be answered in README.md